ChatGPT

Escrowed Conditional Matching for Moral Trade and Moral Public Goods

Executive summary

The highest-confidence conclusion from the three prioritized sources is that a voluntary platform for moral trade should **not** rely on pure goodwill, pure reputation, or pure bilateral bargaining alone. Forethought's analysis argues that moral public goods matter enormously, but that the free-rider problem is "surprisingly difficult to escape" through voluntary contracts at scale; assurance contracts are brittle, dominant assurance helps only a little, and even quadratic funding depends on an external matching pool. Forethought is also pessimistic about social norms as the main engine because they scale poorly and are hard to target well. Ord's paper, by contrast, is optimistic about the *possibility* of moral trade, but emphasizes that practical deployment should start with **simple one-to-one exchanges**, and that the core obstacle is **trust**—especially "counterfactual trust," meaning whether the other party would have acted the same way anyway. He explicitly suggests **escrow** as a partial solution for compliance, while warning that counterfactual trust is much harder to police. [1]

Given that evidence, the single best answer is a **hybrid design centered on escrowed conditional matching**: donors make conditional pledges that only execute if complementary counterparties participate, the platform adds a precommitted matching subsidy to cleared trades, and funds are held in escrow and released only to verified moral public goods. This design combines the feature that is **closest to Ord's conception of moral trade**—conditional exchange across divergent moral preferences—with the feature that has the **strongest empirical record for motivating donations**—matching funds and related price-reduction mechanisms. Matching has repeatedly increased charitable response rates and total revenue in field evidence, while more recent work suggests that "bundling" personally meaningful giving with effective giving can substantially increase effective donations without eliminating donors' emotional attachment to favored causes. [2]

If I had to name only **one mechanism** rather than a full platform architecture, it would be this: **precommitted matching applied only to conditionally cleared, escrowed cross-view donations**. Matching is the donor-facing lever that most reliably changes behavior; conditional clearing makes the donation feel less like unilateral sacrifice and more like mutually beneficial moral trade; escrow and verification solve the trust problem that otherwise destroys scalability. By contrast, pure social signaling is weaker and can crowd out intrinsic motives, pure assurance contracts struggle in large populations, pure bilateral markets suffer from search and verification problems, and pure quadratic-funding systems inherit strong attack surfaces and still require a matching pool. [3]

My confidence in this recommendation is **moderate**. It is high that pure social norms or pure assurance contracts are not the best main mechanism for a general-purpose platform, and moderate that escrowed conditional matching is the best broad design among the mechanisms reviewed. The main reason confidence is not higher is that there is still **very little direct empirical evidence on large-scale moral-**

trade platforms specifically; most of the evidence comes from adjacent literatures on charitable matching, crowdfunding, reputation, lotteries, and public-goods allocation. [4]

What the prioritized sources imply

Moral public goods are a big deal for whether we get a good future

Forethought's first essay frames the problem at the right level: future decision-makers may allocate resources to distant or long-run outcomes that they never personally consume, so the benefits are often moral or ideological rather than personal. That makes **moral public goods** central, and it makes coordination across people with different values unusually important. The essay favors broad power distribution plus strong coordination, and it openly says that current public-goods provision still leans on government taxation because the free-rider problem is severe. [5]

For voluntary mechanisms, the essay is unusually explicit about the limitations. It identifies governments, social norms, and voluntary contracts as the three main candidate approaches, but says voluntary contracts remain highly vulnerable to free-riding even if future technology lowers transaction costs and improves commitment. Forethought also argues that social norms are likely less scalable and less flexible than governments or contracts, and probably cap out at communities of hundreds or thousands. That matters directly for platform design: **reputation and norms may help at the margin, but they are not robust enough to be the core motivational engine for a broad platform.** [6]

The essay's appendix on voluntary contracts is even more decisive. It argues that assurance contracts are brittle when they require universal participation, and still difficult when they use less-than-universal thresholds, because each person cares whether they are pivotal. Forethought says its own modeling suggests that for realistic population sizes—thousands to tens of millions—assurance contracts usually do not succeed unless gains from trade are extremely large. It also notes that quadratic funding relies on an outside source of matching funds, which is itself unexplained in a purely voluntary world. That implies a very practical design principle: **a platform should not bet on pure threshold coordination; it should use thresholding only when paired with a stronger donor incentive, especially subsidized matching.** [7]

Convergence and compromise

The second Forethought essay is less about fundraising mechanics and more about when moral trade is attractive in principle. Its central claim is that there can be very large gains from trade and compromise when groups with different moral views care intensely about different outcomes, especially when they are "resource-compatible" or when shared resources can realize hybrid goods that both groups value. The essay therefore supplies the positive case for a platform: moral trade is not merely compromise under scarcity; it can create **positive-sum allocations**. [8]

But the same essay also sharply limits what a platform should try to do. It states that whether potential gains from trade are actually realized depends in part on whether the **right institutions** exist to support trade. It also emphasizes that bargaining outcomes depend on decision procedures, on power distributions, and on what happens if no agreement occurs. Most importantly, it flags the danger of **threats** and value-destroying extortion. In other words, a platform should facilitate *voluntary Pareto-style exchange*, not arms-race bargaining or harmful side commitments. That strongly favors a platform that only allows **positive,**

pre-vetted funding commitments to registered public goods, and forbids trades that involve harming one cause to benefit another. [9]

The essay also warns that naive collective decision rules can seal off valuable futures. Majority rule can suppress goods that minorities value highly; the same procedure can perform very differently depending on timing; and signaling incentives can distort democratic choice. That implies that a platform should not simply run one-person-one-vote ballots over all funds. If it uses voting or aggregation, it should do so in a way that captures **preference intensity**, filters noise, and protects minorities from being washed out. [10]

Moral Trade by Toby Ord

Ord's paper supplies the closest thing to a direct product blueprint. He gives intuitive examples of moral trade, then moves from informal person-to-person deals to **websites that facilitate matching offers**. He cites vote-pairing platforms such as Nader Trader and VotePair, and political pledge-canceling through Repledge, describing them as "simple markets" that let thousands of people post offers and find matches. He then argues that a mature ecosystem might eventually resemble a moral-trade version of eBay or Craigslist, perhaps with a dedicated currency, but says that early efforts should remain simple and focus on **one-to-one exchanges**. [11]

The most important practical part of Ord's paper is the trust analysis. He distinguishes two kinds of trust. **Factual trust** is whether the other party performs the promised action; **counterfactual trust** is whether they would have done it anyway absent the deal. Ord suggests that escrow can address factual trust by penalizing breach, much as contract law does in ordinary markets. But he emphasizes that counterfactual trust is particularly hard to verify. This is the strongest reason a platform should begin with **observable, short-horizon, directly verifiable commitments**, such as donations to pre-vetted charities or milestone-based grants, instead of lifestyle changes, political gestures, or diffuse behavior-change promises. [12]

Ord's discussion also weakly favors **batch matching over ad hoc negotiation**. His historical examples worked by aggregating many offers and automatically matching them, which reduced search costs. That is important because Forethought's main worry about voluntary coordination is precisely the difficulty of clearing contributions in large populations. A batch-clearing exchange is therefore closer to the spirit of Ord and more realistic than expecting donors to negotiate pairwise manually. [13]

Candidate mechanisms compared

The table below compares the leading candidate mechanisms against the criteria you requested. "IC" refers to incentive compatibility in the rough, practical sense of whether the mechanism induces behavior aligned with what the platform wants while minimizing incentives to misreport, free-ride, or game the system.

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost-effecti
Escrowed conditional matching market	Marries moral trade to a strong donor incentive: donors contribute only if counterparties participate, and receive a lower effective price via matching. Escrow solves compliance trust.	Built from three strong pieces of evidence: matching boosts giving; all-or-nothing/ commitment devices reduce donor risk; hidden-action environments reduce giving and make verification valuable. [14]	**High** for verified donation flows	**High** with batch rounds	**Medium** if identity and audit controls exist	**Medium**; lower if restricted to registered charities	**High**

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost-effecti
Matching funds and micro-matching	Lowers the effective price of giving; makes compromise feel rewarding rather than sacrificial.	Large field experiment found matching increased response and revenue, though larger ratios did not outperform 1:1; recent meta-analysis finds total donations are price-responsive and more responsive under matching than rebates; bundled giving plus matching raised effective giving in Harvard/ Giving Multiplier work. [2]	**Medium-High**	**High** if a match pool exists	**Medium**	**Low-Medium**	**High** v subsid target

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost-effecti
Conditional commitments and assurance contracts	Solve the assurance problem by making contributions contingent on enough others joining.	Forethought is skeptical at large scale; refund-bonus variants improve equilibrium properties and lab performance; all-or-nothing crowdfunding behaves like a useful commitment device. [15]	**Medium**	**Medium** for narrow, thresholded projects; weaker for huge populations	**Medium**	**Medium**	**Mediu** **High** f tightly specif projec

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost-effecti
Bilateral swap markets for moral trade	Closest to Ord's original idea: complementary moral preferences generate gains from trade.	Ord provides credible examples such as vote-pairing and Repledge-style cancellation; direct causal evidence on large-scale welfare gains remains sparse. Vote swapping received First Amendment protection in a Ninth Circuit case, and Repledge-style candidate-charity matching received a favorable FEC advisory opinion in one U.S. context. [16]	**Potentially high** if claims are verifiable	**Low-Medium** unless automated	**High** because of counterfactual trust and strategic bargaining	**Medium-High**, and can be legally sensitive	**Uncer** poten very h niches

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost- effecti
Reputation systems and social signaling	Donors receive prestige, image benefits, or social proof; others may imitate visible giving.	Public recognition increased donation probability by 2.7 percentage points in a Yale field study; social information can increase contributions by about 12% in the best condition; more recent work finds heterogeneity, and image incentives can crowd out intrinsic motivation. [17]	**Low-Medium**	**Medium**	**Medium-High**	**Low**	**Mediu**

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost- effecti
Escrow, verification, and auditing	Increase donor trust by reducing hidden action and misuse.	Ord explicitly recommends escrow for compliance; hidden-action experiments show moral hazard depresses donations and recipient welfare; ex post information can outperform incomplete charitable promises in lab settings. [18]	**High** as an enabling layer	**Medium**	**Lowers other mechanisms' risk**	**Medium-High**	**High** a infrast not as motiva itself

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost-effecti
Voting and aggregation mechanisms	Allocate pooled funds across contested public goods while capturing preference intensity.	Participatory budgeting improves information flows and shifts spending toward user preferences; quadratic voting has strong theory; retro-funding operators now recommend a mix of metrics and expertise because popularity contests and pure project-based voting performed poorly. [19]	**Medium**	**Medium-High**	**Medium-High**	**High**	**Mediu**

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost-effecti
Quadratic funding and related algorithmic matching	Uses many small donations as a signal of broad value and allocates a shared match pool accordingly.	The theory paper claims near-first-best provision under a standard model; Gitcoin documents strong advantages but also sybil and collusion problems; Forethought notes the unresolved dependence on outside matching funds. [20]	**Medium-High** with strong identity	**High**	**High** without robust proof-of-personhood	**High**	**Mediu High** i matur ecosys

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost-effecti
Lotteries and raffle matches	Convert some altruism into excitement or risk-seeking behavior; can stimulate participation.	Theory and experiments show lotteries can outperform simple voluntary contribution mechanisms; a field experiment found revenue gains on both participation and volume margins; recent work suggests some lottery matches can approximate certain 1:1 matches at lower expected cost. [21]	**Medium**	**Medium**	**Medium**	**High**, with substantial charitable-gaming regulation	**Mediu**

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost-effecti
Subscription and club-good models	Use recurring commitments or excludable benefits to smooth funding.	Club theory explains when exclusion can finance quasi-public goods; recurring giving improves retention and predictability, but recurring-charge regulation has tightened and pure club logic fits poorly for genuinely non-excludable moral public goods. [22]	**Medium**	**High**	**Medium**	**Low-Medium**	**Mediu** for acquis better retent

Mechanism	Theoretical rationale	Empirical evidence	IC	Scalability	Manipulation risk	Admin and legal burden	Qualit cost-effecti
Impact certificates and retroactive funding	Fund verified impact after the fact; tradable claims can create stronger ex post evaluation incentives.	Hypercert advocates argue that retro funding improves incentives by rewarding demonstrated impact and building evaluator reputation, but Optimism's own retrospective notes say it still lacks evidence that retro funding produces superior outcomes and recommends mixing metrics with experts. [23]	**Medium**	**Medium-High**	**Medium-High**	**High**	**Mediu** today

The comparative picture is clear. If the question is **what most effectively motivates donors**, the strongest evidence cluster sits around **matching**. If the question is **what most faithfully instantiates moral trade**, the strongest conceptual fit sits around **conditional exchange and batch-matched swaps**. If the question is **what most directly addresses Ord's practical concern**, the answer is **escrow and verification**. The best design therefore combines those three rather than choosing one in isolation. [24]

Why the recommended mechanism wins

The recommended design is an **escrowed conditional matching market**. Each donor creates a pledge that says, in effect: "I will fund this moral public good *if* enough counterparties fund morally different public goods through the same round, and if the platform's precommitted matching pool tops up the cleared trade." A batch-clearing algorithm then matches complementary pledges subject to donor constraints, divides the shared matching subsidy across cleared trades, and sends the money into escrow pending verification. That structure directly operationalizes moral trade, rather than merely running a normal donation platform with a public-goods label. [25]

This beats pure matching because pure matching does not solve the motivational problem at the heart of moral disagreement. Someone who strongly prefers moral good A may still resist donating to moral good B even if B is matched. But a conditional-matching exchange changes the psychological and economic framing: the donor is no longer "abandoning" A for B; they are participating in a **mutually beneficial cross-view bargain** and receiving a subsidy for doing so. That is especially important because real donors often want both personal meaning and effectiveness. The Giving Multiplier evidence is highly relevant here: when donors could split between a favorite charity and a more effective one, donations to the effective option rose sharply, and matching layered on top produced an additional increase. [26]

This also beats pure assurance contracts because Forethought's main objection to voluntary coordination is not just commitment failure but **free-riding at realistic scales**. Matching improves the pivotality calculus by reducing the donor's effective price, and batch clearing reduces search costs. The platform should therefore use conditionality not as a naked threshold game, but as a **screening and alignment tool** layered on top of a stronger price incentive. In practice, that means avoiding "everyone must join" style contracts and instead clearing a large number of partially compatible trades in rounds, with each donor seeing their maximum exposure, minimum accepted counterparty volume, and subsidy schedule in advance. [27]

Most importantly, this design is the cleanest response to Ord's trust problem. Escrow handles **factual trust** because breach or nonperformance can block release. And if the initial platform scope is restricted to **direct transfers to pre-vetted charities or milestone-based grants with objective receipts**, then counterfactual trust becomes much less damaging. The platform should *not* initially support "I promise to become vegetarian" or "I promise not to support policy X" style trades, because those are exactly the domains where Ord says counterfactual trust is hardest and where legal and ethical complexity spikes. A platform that starts with verifiable money flows can still embody moral trade while avoiding the most fragile category of promises. [28]

Finally, this design outperforms pure quadratic-funding or pure retro-funding architectures for launch. Quadratic funding is elegant and may become useful later, but it requires a large matching pool and reliable sybil resistance, and even its defenders document collusion and identity challenges. Retro funding and impact certificates are promising for learning and ex post reward, but even mature operators such as Optimism still say they cannot yet show that retro funding itself produces superior outcomes. These are better viewed as later-stage layers—especially for evaluator-backed allocation of unmatched or retrospective pools—than as the main initial mechanism to motivate first-order donor participation. [29]

Implementation blueprint and governance

Step-by-step implementation

Because target population, budget scale, and jurisdiction were left unspecified, I recommend a **general-purpose voluntary online pilot** aimed at morally diverse donors, but restricted at launch to funding **registered nonprofits and transparently auditable public-goods projects**, not political campaigns or personal behavior-change contracts.

1. **Define the eligible moral public goods registry.**
   Create a pre-vetted registry of eligible causes and recipient organizations. Each entry should include legal status, theory of change, verification standards, and prohibited uses. Use independent evaluators or research partners for diligence, in the style of established evidence-based funders that

gather independent evidence, model cost-effectiveness, review budgets, and follow up on progress. [30]

2. **Segment causes into morally distinct buckets.**
   Donors should be able to tag their pledge as supporting, for example, animal welfare, global health, existential risk reduction, climate, biosecurity, digital welfare, or pluralist charity baskets. But the platform should only count a trade as "moral trade" if the matched counterparty comes from a sufficiently distinct bucket. Otherwise the system degenerates into ordinary matching. This follows Forethought's point that gains from trade are largest when values differ in what they prioritize, yet remain resource-compatible. [31]

3. **Collect conditional pledges with explicit trade parameters.**
   Each donor specifies the amount, acceptable counterpart buckets, a minimum counterparty volume, whether unmatched funds should be refunded or re-routed, and whether they consent to public recognition. This is where the core mechanism lives. The donor is not making an unconditional donation; they are posting a tradeable moral funding intention in a constrained, auditable format. Ord's "simple market" examples suggest this kind of offer-posting and centralized matching is exactly the right starting point. [11]

4. **Attach a precommitted micro-matching schedule.**
   The platform or outside "bridge donors" commit a matching pool in advance. The match should be strongest for donations that clear a cross-bucket trade and, ideally, stronger when the donor allocates some share to a broad compromise basket or high-priority common-good basket. This borrows the central lesson from matching studies and from Giving Multiplier: compromise becomes materially attractive when paired with subsidy. [32]

5. **Run frequent batch-clearing rounds.**
   The platform clears trades every one to two weeks, not continuously. Batch clearing reduces search friction, allows better anti-fraud review, and makes it easier to allocate the match pool fairly. It also matches Ord's historical examples, which were offer-aggregating clearinghouses rather than relationship-based negotiations. [11]

6. **Hold all cleared funds in escrow.**
   Credit-card holds or segregated escrow accounts should be used until the round closes and recipient eligibility is confirmed. For milestone-based grants, release should be staged. This is the platform's practical answer to factual trust. [33]

7. **Verify recipients and publish receipts.**
   The early platform should favor recipients with ordinary charity verification pathways: IRS tax-exempt status where relevant, state registration where required, official payment rails, and standardized reporting. Regulators and donor-protection bodies repeatedly stress charity verification and caution that intermediaries control fund flows and may take fees, so the platform should make every fee, transfer path, and delay explicit. [34]

8. **Use light-touch reputation, not heavy-handed gamification.**
   Public recognition should be optional and privacy-sensitive. The evidence shows some lift from recognition and social information, but also warns about image-crowding and heterogeneity. Offer

badges, verified trade histories, and donor circles as optional features rather than as the core mechanism. [35]

9. **Add a small retrospective pool only after the core mechanism works.**
Once the platform has a record of verified outputs, it can add a small retroactive pool allocated by expert scoring plus metrics. That creates learning incentives without making uncertain retrospective evaluation the main reason donors participate. This follows what retro-funding operators themselves now recommend: combine what data is good at with what experts are good at. [36]

Governance safeguards

The mechanism only works if governance is designed around Forethought's and Ord's failure modes. At minimum, the platform should have:

- **A plural eligibility board** with members from multiple moral viewpoints, rotating terms, and published conflict-of-interest rules. This reduces the chance that one worldview captures the registry. Forethought's concern about power concentration and poor decision procedures makes this non-optional. [37]
- **A hard ban on harmful or extortionary trades.** No commitments to create suffering, spread misinformation, destroy value, or coerce counterparties. This is a direct response to Forethought's warning about threats. [38]
- **Open-source or fully auditable clearing logic** for matching, subsidy allocation, identity thresholds, and dispute handling.
- **Identity and anti-sybil controls** proportionate to match size. Where feasible, use proof-of-personhood, payment-method checks, graph-based anomaly detection, and manual review on suspicious clusters. Quadratic-style systems document how necessary this is. [39]
- **Treasury separation** between operating funds, the matching pool, and recipient disbursement funds.
- **Explicit recurring-payment consent and easy cancellation** if the platform later adds subscriptions. Current U.S. consumer-protection rules emphasize express informed consent and easy cancellation for recurring charges. [40]
- **Fraud reporting and appeal channels** for donors, charities, and counterparties, with escalation to regulators when appropriate. [41]

Metrics for success

The platform should judge itself by **incremental, verified, cross-view funding**, not by raw gross volume alone.

- **Incremental dollars to moral public goods per subsidy dollar.**
- **Trade-clear rate:** share of posted conditional pledges that clear.
- **Counterparty-diversity index:** how often trades occur across morally distinct buckets rather than within one bucket.
- **Verification pass rate and time-to-disbursement.**
- **Donor retention and repeat cross-view participation.** Recurring or repeated donors matter because donor retention is a known platform challenge. [42]
- **Manipulation rate:** sybil flags, reversed allocations, disputes, fraudulent recipients.

- **Additionality estimate:** behaviorally measured share of matched dollars that would not have been donated absent the platform.
- **Donor trust score:** post-round surveys on fairness, clarity, and confidence in fund use.
- **Recipient outcome quality:** verified milestones, evaluator scores, and downstream impact evidence where feasible. [30]

Failure modes and mitigations

The most important expected failures are not mysterious.

- **Free-riding or low clearing rates.**
  Mitigation: use matching subsidies, smaller rounds, and round-specific cause segmentation; allow partial clearing instead of all-or-nothing universal thresholds. [43]

- **Counterfactual dishonesty.**
  Mitigation: begin with direct donations and milestone-based grants, not behavioral-lifestyle promises; require recipient-side verification rather than donor self-report. Ord's analysis strongly supports this scoping choice. [12]

- **Sybil attacks and collusion.**
  Mitigation: verified identity, pooled anomaly detection, contribution caps for subsidy eligibility, cooling-off periods, and ex post adjustment of allocations. [39]

- **Fraud or hidden action by recipients.**
  Mitigation: escrow, staged release, audited receipts, and evaluator review. Hidden-action experiments show this matters to donor behavior and recipient welfare. [44]

- **Popularity contests or ideological capture.**
  Mitigation: do not allocate the main pool by simple voting; where pooled discretionary funds exist, use expert-plus-metrics review and publish rationales. [45]

- **Legal trouble in regulated domains.**
  Mitigation: start with registered charities and standard philanthropic transfers; do not launch with campaign finance, electoral swaps, or regulated gaming unless counsel clears each jurisdiction. Both vote-swapping and candidate-charity matching have some favorable U.S. precedents, but they remain special cases, not universal safe harbors. [46]

- **Reputational damage from scam look-alikes or fake recipients.**
  Mitigation: donor-facing verification checkmarks, official payment routing through registered entities, and strong anti-impersonation messaging. [47]

Pilot timeline

The pilot should start narrow: only direct payments to verified recipients, modest match caps, and batch rounds every two weeks. That timeline follows Ord's advice to start simple and Forethought's skepticism about too-ambitious voluntary coordination schemes. [48]

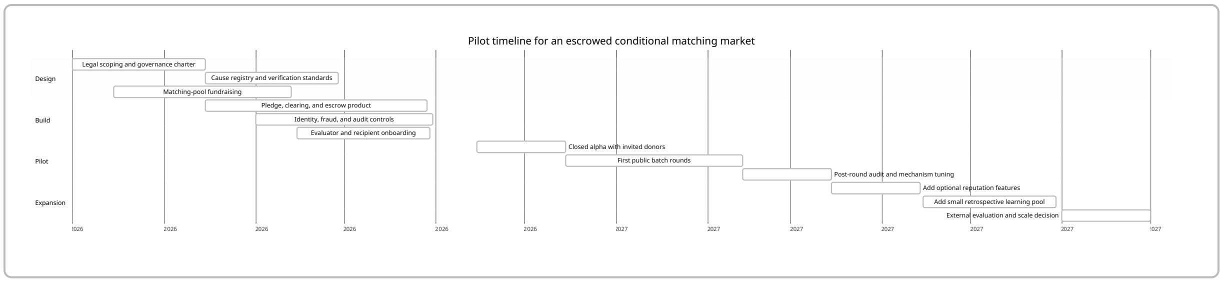

Open questions and limitations

The biggest limitation is that there is **no mature, large-scale, published RCT of an actual moral-trade platform** in the Toby Ord sense. The closest evidence comes from adjacent mechanisms: charitable matching, public-goods crowdfunding, political cancellation platforms, reputation experiments, and retro-funding pilots. That is enough to rank mechanisms with moderate confidence, but not enough to claim precise effect sizes for a full moral-trade platform. [49]

A second limitation is that matching depends on the availability of a **bridge-donor subsidy pool**. Forethought is right that purely voluntary settings do not automatically explain where this pool comes from. In practice, the answer is likely a separate class of pluralist or impact-focused funders who value creating trades that unlock otherwise unreachable donations. But that is still a dependency, not a solved theorem. [50]

A third limitation is jurisdiction. Electoral swaps, raffles, and some impact-claim markets face materially different legal treatment across places. The recommendation here is therefore strongest for **charity-to-charity or donor-to-charity funding flows** and intentionally weaker for political or quasi-security instruments. [51]

Core sources

The three prioritized sources were the decisive anchors for this recommendation: Forethought's **"Moral public goods are a big deal for whether we get a good future"**, which diagnoses the free-rider and scalability problem for voluntary coordination; Forethought's **"Convergence and Compromise"**, which explains when moral trade creates large positive-sum gains and why institutions and anti-threat safeguards matter; and Toby Ord's **"Moral Trade"**, which points toward simple matching markets, one-to-one exchanges, and escrow as the practical starting point. [52]

The most important additional sources were Dean Karlan and John List on donation matching, the recent meta-analysis of price elasticity under rebates and matches, the Giving Multiplier research and platform evidence on donation bundling plus matching, the provision-point and crowdfunding literature on conditional commitments, Ariely–Bracha–Meier and Shang–Croson on image and social information, Buterin–Hitzig–Weyl on quadratic funding, Gitcoin and Optimism materials on real-world algorithmic public-goods funding, and official FTC/FEC/BBB/GiveWell materials on verification, recurring-payment compliance, and recipient due diligence. [53]

[1] [4] [5] [6] [7] [15] [27] [43] [48] [50] [52] Moral public goods are a big deal for whether we get a good future
https://www.forethought.org/research/moral-public-goods-are-a-big-deal-for-whether-we-get-a-good-future

[2] [14] [24] [32] [49] [53] Does Price Matter in Charitable Giving? Evidence from a Large-Scale Natural Field Experiment - American Economic Association
https://www.aeaweb.org/articles?id=10.1257%2Faer.97.5.1774

[3] Doing Good or Doing Well? Image Motivation and Monetary Incentives in Behaving Prosocially - American Economic Association
https://www.aeaweb.org/articles?id=10.1257%2Faer.99.1.544

[8] [9] [10] [31] [37] [38] Convergence and Compromise: Will Society Aim for Good Futures?
https://www.forethought.org/research/convergence-and-compromise

[11] [12] [13] [16] [18] [25] [28] [33] Moral Trade
https://amirrorclear.net/files/moral-trade.pdf

[17] [35] Public Recognition and Fundraising in the United States | The Abdul Latif Jameel Poverty Action Lab
https://www.povertyactionlab.org/evaluation/public-recognition-and-fundraising-united-states

[19] Participatory budgeting and the pattern of local government ...
https://www.sciencedirect.com/science/article/abs/pii/S0176268022000465?utm_source=chatgpt.com

[20] [29] A Flexible Design for Funding Public Goods | Management Science
https://pubsonline.informs.org/doi/10.1287/mnsc.2019.3337

[21] econweb.ucsd.edu
https://econweb.ucsd.edu/~jandreon/PhilanthropyAndFundraising/Volume%202/2%20Morgan%202000.pdf

[22] An Economic Theory of Clubs
https://www.jstor.org/stable/2552442?utm_source=chatgpt.com

[23] Hypercerts: A new primitive for public goods funding | Protocol Labs
https://www.protocol.ai/blog/hypercert-new-primitive/

[26] Making Charitable Giving More Competent | Harvard Magazine
https://www.harvardmagazine.com/2023/04/right-now-charitable-giving

[30] GiveWell | Charity Reviews and Research
https://www.givewell.org/

[34] Donating Through Crowdfunding and Fundraising Platforms | Consumer Advice
https://consumer.ftc.gov/articles/donating-through-crowdfunding-and-fundraising-platforms

[36] [45] Lessons learned from two years of Retroactive Public Goods Funding - Retro Funding Missions - Optimism Collective
https://gov.optimism.io/t/lessons-learned-from-two-years-of-retroactive-public-goods-funding/9239

[39] Quadratic Funding | Gitcoin
https://gitcoin.co/mechanisms/quadratic-funding

[40] Federal Trade Commission Announces Final “Click-to-Cancel” Rule Making It Easier for Consumers to End Recurring Subscriptions and Memberships | Federal Trade Commission
https://www.ftc.gov/news-events/news/press-releases/2024/10/federal-trade-commission-announces-final-click-cancel-rule-making-it-easier-consumers-end-recurring

[41] [47] Wise Giving Wednesday: How to Spot Charity Impersonation Scams - Give.org
https://give.org/news/charity-impersonation-scams

[42] Donor retention in online charitable crowdfunding platform
https://www.sciencedirect.com/science/article/abs/pii/S0167923620301822?utm_source=chatgpt.com

[44] Charity Fraud: An Experimental Study of the Moral Hazard Problem in the Charity Market
https://hias.hit-u.ac.jp/wp-content/uploads/2024/04/HIAS-E-139.pdf

[46] cdn.ca9.uscourts.gov
https://cdn.ca9.uscourts.gov/datastore/opinions/2007/08/06/0655517.pdf

[51] AO 2015-08: Company May Operate Candidate Contribution Charitable Match Platform
https://www.fec.gov/updates/ao-2015-08-company-may-operate-candidate-contribution-charitable-match-platform/