

Improving Moral Trade Funding for Moral Public Goods

Executive summary

The short answer is **no**: the current mechanism on Moral Trade is **not** the most effective possible mechanism for a voluntary platform that aims to motivate spending on moral public goods. It is, however, **directionally correct**. Moral Trade has already moved toward the right design family by focusing on (i) shared public goods rather than only bilateral barter, (ii) thresholded commitments, (iii) identity and review gates, (iv) sponsor matching, and (v) capped quadratic-style bonus allocation after campaigns clear thresholds. Those choices fit the core lessons of Toby Ord's paper on moral trade, Forethought's note on moral public goods, and Forethought's broader discussion of convergence, compromise, trade, and threats.

1

But the current implementation still looks like a **prototype of the right idea**, not the strongest feasible execution of it. Moral Trade's own public pages say the platform is a pilot, has **0 completed agreements, 0 live offers** on the homepage snapshot, treats public-goods coordination as "mechanism research," has **integrated checkout planned but not active**, reports **0% auto-verified share of contributions**, and shows **\$0 released in ledger** for the inspected demo round even though it publishes payout plans and threshold-cleared campaigns. Those are clear signs that the current mechanism is not yet extracting the full available gains from trade or solving the friction problems that matter most in practice.

2

The strongest feasible upgrade for **moraltrade.org as a voluntary platform** is a **hybrid mechanism** I will call **Verified Coalition Quadratic Assurance with Failure Bonuses**. In user-facing language, this should be presented as a **Common Ground Budget**: users set a recurring budget for public goods they see as strong priorities or acceptable common-ground goods; the platform then routes those pre-authorized budgets across thresholded campaigns, unlocks sponsor-funded base matching only after review and threshold gates pass, and allocates bonus matching using an identity-weighted, diversity-aware quadratic formula. Small, sponsor-funded failure bonuses should be added for early contributors when a campaign fails, because refund-bonus literature suggests that such bonuses improve threshold crowdfunding, while pure one-shot assurance contracts remain brittle.

3

My confidence is **high** that the current mechanism is not the most effective feasible one for Moral Trade in its present form, because the site's own metrics and product boundaries show unresolved friction and non-live behavior. My confidence is **moderate** that the hybrid I propose is the best practical upgrade under current voluntary-platform constraints, because Forethought is probably right that fully solving moral-public-goods free riding with voluntary contracts is intrinsically hard; if coercive taxation or acausal coordination were available, those could outperform any ordinary web platform.

4

What the prioritized sources show

Toby Ord's paper makes three points that matter most here. First, moral trade can generate large gains when people care intensely about different moral outcomes, but there are “almost no mechanisms” for facilitating it today, and a mature ecosystem would likely require something more like **markets, bargaining, and perhaps even a moral-trade currency**, rather than only ad hoc pairwise swaps. Second, the biggest practical obstacle is **trust**, especially both **factual trust** and **counterfactual trust**: did the action really happen, and would it have happened anyway? Third, moral trade can create **perverse incentives** and negative externalities unless baselines, enforcement, and review are handled carefully. Those are exactly the design problems any serious platform must optimize around. ⁵

Forethought's note on moral public goods sharpens the mechanism-design picture. Its core claim is that when many people share **some** consensus preferences while still differing strongly on private or idiosyncratic preferences, coordination around moral public goods can create extremely large gains. In their toy model, with distributed power and successful coordination, each dollar given up in idiosyncratic goods effectively buys **N dollars** of consensus goods because everyone else does the same; the note explicitly concludes that we should want both **widely distributed power** and **good coordination mechanisms** for shared goods. Just as importantly, it argues that standard one-shot voluntary mechanisms are much weaker than they may first appear: assurance contracts are brittle, dominant assurance contracts help only somewhat, and even with advanced technology the free-rider problem is hard to escape. It also notes that quadratic funding requires an **outside matching pool**, not magic. ⁶

Forethought's broader essay on convergence and compromise adds one more crucial lesson: **not all moral trade scales equally well**. Trade over “hybrid goods” can sometimes produce gains, but for many linear, resource-hungry value systems the gains from compromise may be limited unless the right institutions exist. That makes moral public goods especially attractive as a scalable wedge: they can be valued “somewhat” by many people without requiring nearly optimal hybrid outcomes for each worldview. The essay also emphasizes that trade only delivers a mostly-good future if the institutions are good enough, if mutual gains are actually realized, and if the system avoids threats and extortion dynamics. ⁷

Taken together, the prioritized sources imply a fairly crisp design target for Moral Trade: **do not center the product on bilateral pair matching alone**; instead, use the platform to discover overlapping support, convert that overlap into conditional pooled funding, make trust legible through verification and review, and guard hard against anti-threat violations and baseline gaming. That is very close to what Moral Trade is trying to do, but its present implementation does not yet fully convert weak common-ground support into actual money flows. ⁸

What moraltrade.org currently does

Moral Trade presents itself as “a pilot institution for cooperation under disagreement,” not as a mature marketplace. The homepage says users can start with a **pledge swap, donation offset, or public-good commitment**, while keeping baselines, evidence rules, and manual review explicit. The same page reports **0 live offers, 8 worked examples, 2 public profiles, and 0 completed agreements**, and it emphasizes “no custody or escrow,” “manual review before reliance,” and privacy-first matching. That is a safety-conscious posture, but it is also a strong signal of low demonstrated liquidity and low validated throughput. ⁹

The public-goods mechanism is the **Public Goods Fund**, also called the **Moral Public Goods Fund** or MPGF. Moral Trade says this is the platform’s “main test” of moral-public-goods coordination. The public page describes the mechanism as a combination of **verified contribution intents, conditional payment authorization, sponsor matching, dissent notes, and reviewer verification**, operating through threshold commitments rather than immediate uncapped donations. It explicitly says its “motivating layer is verified assurance matching,” and the round page further labels the inspected round as **Verified Quadratic Assurance Funding**. The round combines a sponsor pool, threshold gates, base matching, and a bonus layer allocated after threshold and review conditions pass. ¹⁰

The current round mechanics are unusually transparent. For the inspected May 2026 demo round, the site publishes a **\$1,500 sponsor pool**, a **\$1,000 base match budget**, a **\$500 QF bonus budget**, a **\$100 counted cap per donor**, **9 verified donors**, and an identity policy in which **identity confidence can affect eligibility or QF weight but “moral reputation never affects allocation power.”** It also publishes campaign-level thresholds, donor counts, common-ground scores, and final match estimates. That is a real strength: the allocation logic is legible, donor power is not tied to moral status, and identity is being used mainly as a Sybil defense. ¹¹

Moral Trade also already has a **common-ground discovery layer**. The MPGF candidate-pools page says campaign ordering ranks **“coordinatability, not moral truth,”** using private-by-default support signals, cross-cluster breadth, threshold progress, and reviewability. The page exposes “common-ground scores,” distinguishes **consensus** from **hybrid** goods, and lets users think in terms of broad moral clusters such as humanitarian, longtermist, animal-inclusive, and institutional-pluralist lenses. This is directly aligned with the theory that moral public goods are goods many moral views can value somewhat. ¹²

The trust and safety model is also unusually explicit. Moral Trade separates **action evidence, baseline confidence, and externality review**; it exposes structured review states; it blocks anti-threat proposals; and it publishes factor codes, appeals, privacy rules, and review scopes. Its validation page draws a clean line between narrow evidence claims and broader claims of moral value, legal enforceability, tax treatment, or outcome guarantees. The safety and primer materials also emphasize a no-trade baseline, prohibited coercive proposals, privacy-safe previews, no autonomous outreach, and reviewer-bounded decision-making. These are the right norms for Ord’s trust problem and for Forethought’s warnings about threats. ¹³

That said, the public metrics show several obvious weaknesses. The funding-metrics page reports a **50% threshold-clear rate, 200% sponsor leverage, 0% auto-verified share of contributions, and 238 hours from pledge to counted contribution**. The round page shows **\$0 released in ledger** despite publishing a total payout plan, and the main public-goods page says integrated checkout is **planned, not active**. The transparency report also shows many zero counts across reviewed suggestions, intro packets, appeals, evidence reviews, and unresolved disputes, which is consistent with a pilot that has not yet accumulated enough operational data to validate broad-market behavior. ¹⁴

Is the current mechanism the most effective possible

My bottom-line judgment is: **no, but it is close to the right design family**. The current mechanism earns high marks for theoretical alignment and for being unusually self-aware about review, identity, privacy, baseline integrity, and anti-threat rules. It already recognizes that public-goods coordination is the scalable extension of bilateral moral trade; it already uses sponsor matching and capped QF; it already distinguishes

common-ground discovery from moral ranking; and it already routes through verification and human review. On the theory side, those are exactly the ingredients one would expect from the prioritized sources.

15

The main problem is that the current mechanism still behaves too much like a **static campaign-and-threshold interface** rather than a **market that actively converts partial overlap into funding**. Forethought's core point about moral public goods is that many people may value a shared good **somewhat**, even when almost none would choose it as their top spending priority. Current Moral Trade surfaces "weak common-ground support" and cluster breadth, but it does not yet appear to turn that information into a first-class budgeting primitive for donors. In other words, weak overlap is being used mostly for ranking and explanation, not yet as the main engine of allocation. 16

That matters because assurance-style systems are brittle. Forethought explicitly models assurance contracts and says they "usually won't go through" for realistic population sizes unless gains from trade are extremely large. Their critique of dominant assurance contracts is also important: a failure payment can help a bit, but it does not remove the deeper incentive to free ride when others may fund the public good anyway. So a mechanism that still leans primarily on users explicitly choosing one campaign and one thresholded pledge is likely leaving a lot of weak common-ground support uncollected. 17

The public-goods literature points toward a better hybrid. Quadratic funding is attractive because, with a sponsor pool, it amplifies broad support rather than only large donors. But QF works best when there is a philanthropic or public matching pool and when Sybil/collusion risk is controlled. Moral Trade already has the beginnings of that infrastructure: sponsor-pool refill sources, identity-weighted eligibility, donor caps, and public review. Meanwhile, assurance/refund-bonus literature suggests that failure bonuses can reduce first-mover hesitation and improve threshold crowdfunding. The strongest mechanism is therefore not "pure assurance" or "pure QF," but a **coalition-forming hybrid**: provision points to avoid paying into doomed campaigns, sponsor-funded base match to make passage salient, and diversity-aware QF to reward broad, cross-view support once passage is feasible. 18

Here is the synthesis in compact form.

Criterion	Current mechanism	Proposed mechanism	Judgment
Incentive compatibility	Better than plain donations because of thresholds, review, and matching, but still relies too much on campaign-by-campaign donor initiative	Converts recurring "weak common-ground" willingness into actual contingent budgets; adds small failure bonuses to reduce hesitation	Proposed is stronger
Scalability	Better than bilateral barter, but still brittle and operationally manual	Coalition-based routing, recurring budgets, and automated round clearing scale much better	Proposed is much stronger

Criterion	Current mechanism	Proposed mechanism	Judgment
Robustness to manipulation	Strong anti-threat, review, and identity posture; QF still needs careful anti-Sybil handling	Keeps current review posture while making identity, caps, and diversity constraints explicit in allocation math	Proposed is somewhat stronger
Fairness	Good separation of moral reputation from allocation power	Preserves that rule and adds cross-cluster breadth requirements so one worldview cannot easily dominate “common-ground” pools	Proposed is stronger
Cost-effectiveness	Sponsor match and QF are good; 0% auto-verification and long review latency waste attention and money	Uses sponsor funds where they matter most, routes only threshold-feasible budgets, and reduces dead-end user effort	Proposed is stronger
User experience	Safe but high-friction; still demo-heavy; integrated checkout inactive	One recurring budget, clearer project stances, less repeated choice friction, better near-threshold nudges	Proposed is much stronger
Measurability	Good public metrics but little live data and many zeroes	Adds direct measurement of weak-support conversion, coalition breadth, and payout lag	Proposed is stronger
Alignment with moral-public-goods theory	High	Even higher, because it directly operationalizes “many people value somewhat” rather than merely exposing it	Proposed is strongest

These ratings synthesize the cited theory and the site’s current public metrics and flows. 19

Recommended mechanism

The single best upgrade is a **hybrid** I would deploy under a user-facing label like **Common Ground Budget**, with the technical name **Verified Coalition Quadratic Assurance with Failure Bonuses**.

The mechanism works like this. Each user sets a recurring monthly budget for moral public goods and marks each pool as one of three stances: **strong support**, **weak common-ground support**, or **dissent / review needed**. The platform then uses only the first two stances as allocatable budget, subject to user caps and opt-outs. Instead of making each donor manually bet on a single threshold campaign, the system routes approved budget across campaigns that are (i) review-approved, (ii) threshold-feasible, and (iii) supported by a sufficiently broad coalition of moral clusters. Budget is only captured after **identity, review, threshold, and challenge** gates pass. Sponsor money then enters in two layers: a **base assurance match** that only unlocks after a campaign clears hard thresholds, and a **bonus layer** using a capped, identity-weighted, diversity-aware quadratic formula to amplify broad support. When a campaign fails, early

contributors can receive a **small sponsor-funded failure bonus** or carry-forward credit, enough to reduce hesitation but not so large that the platform becomes a failure-lottery. ²⁰

Why this mechanism, specifically?

It preserves what Moral Trade already gets right: **no hidden moral ranking**, no use of “moral reputation” in allocation power, strong anti-threat rules, identity-aware Sybil resistance, public review, and nontrivial matching. It also respects the strongest lesson from Ord: the platform should not stop at moral barter; it should build more market-like institutions that reduce search and contracting costs while still handling trust and counterfactual-baseline problems. At the same time, it absorbs Forethought’s critique of voluntary mechanisms by **not** betting everything on a single unanimous assurance contract. Instead, it lets subsets of users with partial overlap actually coordinate as a **coalition**, which is much closer to Forethought’s own claim that even when convergence is limited, any subset with shared consensus goals can still gain from trade. ²¹

The biggest conceptual change is this: **weak common-ground support must become spendable budget, not just a discovery signal**. That is the main gap in the current site. Right now, Moral Trade can tell a user that a campaign has cross-cluster breadth or weak common-ground support. The improved mechanism should let that user say, in effect, “I would not personally top-fund this, but I am willing to put up to \$20/month behind it if enough other verified people from other moral clusters do the same.” That is exactly the kind of motivation moral public goods theory predicts. ²²

One more nuance matters. If the comparison class is **all possible mechanisms in the world**, then Moral Trade cannot beat a competent government with legitimate taxation power, and it cannot implement Evidential Cooperation in Large Worlds in any ordinary product sense. Forethought explicitly thinks democratic taxation can outperform voluntary deals for causal coordination, and ECL could be even stronger in some acausal settings. So the claim here is narrower and more practical: **for a voluntary platform with Moral Trade’s current posture, the best-supported build target is the coalition-based assurance-plus-matching hybrid above**. ²³

Codex build instructions

The instructions below are written for **Codex GPT-5.5 xhigh reasoning** as an implementation brief for moraltrade.org.

Build objective

Replace the current “campaign-by-campaign pledge” public-goods flow with a **round-based Common Ground Budget system** that:

1. collects recurring user budgets and explicit support stances,
2. routes only approved budget into threshold-feasible campaigns,
3. unlocks sponsor-funded base match only after hard gates pass,
4. allocates bonus funds using capped, identity-weighted, diversity-aware quadratic matching,
5. preserves current privacy, review, anti-threat, and no-moral-ranking guarantees, and

6. materially lowers contribution friction and time-to-counted-contribution. This is the core product objective because the current public metrics show friction remains high and auto-verification remains absent. ²⁴

Preserve these current invariants

Do **not** regress the following current design constraints:

- No global moral ranking.
- Moral reputation must not change allocation power.
- Anti-threat baseline rules remain a blocking gate.
- Action evidence, baseline confidence, and externality review stay separate.
- Public pages remain privacy-safe and aggregate-first.
- Matching and suggestion systems remain deterministic and review-bounded for reliance-bearing states unless a new governance contract explicitly broadens this. ²⁵

Extend the existing data model

Moral Trade already exposes a validator-backed data model with participants, profiles, offers, evidence claims, artifacts, payment records, appeals, disputes, privacy grants, match suggestions, and related provenance objects. Extend that model rather than replacing it. ²⁶

Add these entities:

```
mpgf_round
- id
- slug
- status {draft, open, locked, reviewing, payable, released, closed}
- opens_at
- closes_at
- review_freeze_at
- base_match_budget_cents
- bonus_budget_cents
- failure_bonus_budget_cents
- donor_counted_cap_cents
- min_verified_supporters_default
- min_active_clusters_default
- qf_formula_version
- diversity_formula_version
- threshold_policy_version
- sponsor_pool_source_hash
- calculation_hash

public_good_project
- id
- round_id
- pool_id
```

- title
- cause_area
- good_type {consensus, hybrid}
- summary
- destination_type {external_charity, fiscal_host, signed_sponsor_route}
- destination_reference
- requested_max_cents
- minimum_viable_cents
- threshold_amount_cents
- threshold_supporters_min
- threshold_clusters_min
- review_state
- challenge_state
- dissent_state
- base_match_ratio_bps
- qf_cap_multiple_bps
- milestone_schedule_json

common_ground_budget

- id
- participant_id
- round_id
- monthly_budget_cents
- payment_method_ref
- budget_state {draft, active, paused, expired}
- default_visibility
- fallback_rule {carry_forward, re-route, release_hold}
- created_at
- updated_at

project_support_stance

- id
- participant_id
- project_id
- stance {strong, weak, dissent, abstain}
- max_alloc_cents
- max_alloc_pct
- rank_order
- cluster_tag_snapshot
- note_redacted
- created_at

budget_allocation_intent

- id
- round_id
- participant_id
- project_id
- alloc_cap_cents

- router_weight_bps
- counted_cap_cents
- eligibility_state
- authorization_state
- capture_state
- review_state
- appeal_state

identity_eligibility

- participant_id
- round_id
- human_verified_bool
- identity_confidence_score
- sybil_flag_state
- counted_weight_bps
- reviewed_at

round_allocation_result

- id
- round_id
- project_id
- direct_counted_cents
- base_match_cents
- bonus_match_cents
- failure_bonus_cents
- final_payout_plan_cents
- active_supporters_count
- active_clusters_count
- qf_score_raw
- qf_score_adjusted
- diversity_factor
- calculation_hash

failure_bonus_claim

- id
- round_id
- participant_id
- project_id
- qualified_bool
- bonus_cents
- state {pending, granted, expired}

Keep all provenance, artifact, challenge, and reviewer-decision links consistent with the existing validator model. Every allocation result should be reproducible from an append-only calculation input bundle. 27

Implement the allocation algorithm

Use a **hard-gates-first, bonus-second** pipeline.

Hard gates

A project becomes **threshold-cleared** only if all of the following are true:

- review state is approval-compatible,
- challenge window has closed or is explicitly non-blocking,
- direct counted budget meets threshold amount,
- counted verified supporters meet the supporter minimum,
- active moral clusters meet the cluster minimum,
- no anti-threat or unresolved destination-proof blocker is open. 28

Direct counted budget

For each user-project pair, compute:

```
counted_contribution =  
min(  
    user_project_alloc_cap,  
    per_donor_counted_cap,  
    user_remaining_round_budget,  
    project_remaining_threshold_gap_or_postclear_cap  
) * identity_weight
```

Where:

- `identity_weight` depends only on identity confidence / eligibility policy,
- **not** on cause preference, ideology, reputation, or support-signal strength beyond the user's own explicit stance. That matches the current site's published policy. 11

Stance weights

Use:

- `strong = 1.0`
- `weak = 0.6`
- `dissent = 0.0` for allocation, but increments review pressure
- `abstain = 0.0`

The exact weights are not sacred. They should be configuration, not hard-coded constants.

Base match

After hard gates pass, apply:

```
base_match_cents =
min(
    round_base_budget_remaining,
    direct_counted_cents * base_match_ratio
)
```

Default initial `base_match_ratio = 1.0` for continuity with the current demo. ¹¹

Bonus match

For threshold-cleared projects, compute a capped QF-style score:

```
qf_raw(project) = (Σ_u sqrt(counted_contribution_u * stance_weight_u))^2 - Σ_u
counted_contribution_u
```

Then adjust for coalition diversity:

```
diversity_factor =
1 + alpha * normalized_cluster_diversity
    - beta * normalized_dissent_pressure
```

with a hard cap, for example `[0.75, 1.25]`.

Then:

```
qf_adjusted(project) = max(0, qf_raw(project) * diversity_factor)
bonus_match(project) = bonus_budget * qf_adjusted(project) / Σ_all_cleared
qf_adjusted
```

This is the place where Moral Trade's common-ground thesis becomes actual money movement rather than directory ranking. Keep donor caps and identity filters in place to reduce Sybil/QF gaming, because real QF deployments treat Sybil resistance as a central operational requirement. ²⁹

Failure bonus

Add a **small** failure-bonus pool for early contributors to campaigns that did not clear thresholds, with bonuses proportional to early qualified counted intent and capped tightly per user. Keep this secondary. The design rationale is that refund-bonus mechanisms can improve threshold crowdfunding and reduce hesitation, but Forethought's critique means the platform should not pretend this alone solves free-riding.

³⁰

A simple version:

```
qualified = project_failed && user_authorized_before_early_cutoff &&
no_sybil_flag
failure_bonus_cents =
min(per_user_bonus_cap, floor(failure_bonus_pool * user_failed_counted /
total_failed_qualified_counted))
```

Redesign the public UX

The new UX should have three primary screens.

Common Ground Budget setup

```
-----
Common Ground Budget
Set one monthly budget for public goods you can endorse.

Monthly budget: [$50]
Capture rule: [only after threshold + review pass]
Fallback if project fails:
(o) re-route to next eligible common-ground project
( ) carry forward to next round
( ) release hold

Your stance on current pools:
[Strong] Global health and basic needs
[Weak common-ground] Existential-risk resilience
[Weak common-ground] Public-interest knowledge
[Dissent / review] Animal welfare transition

Per-project max:
Global health      [$25]
X-risk             [$15]
Knowledge          [$10]
Animal welfare     [$0]
-----
```

This screen is the core product change. The goal is to let users convert partial overlap into pre-committed budget without forcing them to “top-rank” every pool.

Round board

```
-----
May Round | Common Ground Board
```

Global health and basic needs
Threshold: 3 supporters + \$250 + 2 clusters
Status: Cleared
Your stance: Strong
Your max allocation: \$25
Projected direct counted: \$23.50
Projected match: \$49.97

Existential-risk resilience
Threshold: 3 supporters + \$500 + 2 clusters
Status: Near threshold
Your stance: Weak common-ground
Add \$5 to help clear? [Yes]

Public-interest knowledge
Threshold: 2 supporters + \$180 + 2 clusters
Status: Needs one more supporter
Your stance: Weak common-ground

Surface threshold proximity and whether a user's weak support is currently pivotal. That is where the mechanism should create motivation.

Contribution state and proof ledger

Your Common Ground Budget | Round status

Authorized budget: \$50
Currently routed:
- Global health: \$23.50 counted
- X-risk: \$15 pending threshold
- Knowledge: \$0 failed threshold
Failure bonus earned: \$1.20

Review states:
- Identity: passed
- Threshold: mixed
- Destination proof: passed for GH
- Challenge window: closed for GH, open for X-risk

Payout milestones:
- GH milestone 1: pending
- GH milestone 2: pending

This screen should preserve the current transparency ethos while making the mental model much simpler.

Implement payment flow, not smart contracts, in v1

Do **not** use on-chain smart contracts for the first serious implementation. Moral Trade's current mechanism depends on human review, privacy controls, redaction, and identity-verified donor uniqueness. That is much better served by ordinary payment infrastructure with strong audit logs than by premature on-chain exposure. The correct v1 payment flow is a **saved payment method plus just-in-time authorization/capture** flow. ³¹

Use Stripe's official primitives as follows:

1. Identity

Use Stripe Identity for KYC-grade identity verification where legally appropriate. Stripe documents support document and selfie checks, and Moral Trade's current product direction already depends on identity confidence as a QF-weight/eligibility control. ³²

2. Payment method collection

Save a reusable payment method up front. Do **not** rely on a long-lived authorization across a whole monthly round, because Stripe manual-capture authorizations expire after a limited window. ³³

3. Just-in-time authorization

Near round close, for threshold-feasible contributions only, create PaymentIntents or equivalent authorizations and confirm them close enough to the review/capture window that they will not expire before capture. Stripe explicitly supports manual capture and webhook-driven event handling. ³⁴

4. Capture after review

Capture only after the project clears threshold, review, and challenge gates. If a project fails or becomes blocked, cancel or release the authorization and, if applicable, post a small failure bonus from the sponsor pool. ³⁵

5. Fallback evidence path

Preserve the site's current Every.org/manual-evidence fallback until direct partner-legal readiness is complete. The governance page already distinguishes partner/fiscal-sponsor custody from platform custody; keep that separation. ³⁶

Ship the reviewer and governance changes

Codex should also implement these governance-side product changes:

1. Add a **Coalition Review Console** that shows:
2. project thresholds,
3. active supporter count,

4. active cluster count,
5. dissent pressure,
6. destination proof state,
7. challenge-window state,

8. allocation simulation preview.

9. Add a **Threshold Calibration** panel for rounds, because the site already publishes a threshold-calibration artifact in governance and should turn that into a real operator workflow. ³⁷

10. Add a **Sponsor Pool Flywheel** dashboard that explicitly shows refill sources:

11. direct sponsor deposits,

12. recurring member tithes,

13. donation-offset surplus,

14. trade-surplus tithes.

This already exists conceptually on the governance page and should become a first-class operational object. ³⁸

15. Add a **public postmortem template** per round:

16. weak-support conversion rate,

17. threshold-clear rate,

18. donor retention,

19. appeal rate,

20. Sybil flag rate,

21. payout lag,

22. sponsor leverage,

23. bonus-vs-base allocation split.

The platform already publishes postmortem and KPI artifacts; the proposed mechanism needs additional coalition-specific metrics. ³⁹

Measure these KPIs

Keep the current aggregate-only, privacy-safe measurement posture, but add mechanism-specific KPIs.

Core funding KPIs:

- weak-support-to-counted-dollar conversion,
- strong-support-to-counted-dollar conversion,
- threshold-clear rate,
- average active clusters per cleared project,
- sponsor leverage ratio,
- failure-bonus utilization rate,

- time from budget activation to counted contribution,
- time from counted contribution to payout release,
- donor retention into next round,
- share of projects with at least one dissent-triggered review. 40

Safety KPIs:

- duplicate identity flags,
- reviewer overturns,
- blocked-project precision,
- anti-threat false positives,
- challenge-window reopen rate,
- privacy leakage incidents = target zero. 41

Run these A/B tests

Run four disciplined A/B tests, sequentially, not all at once.

Test one

Current per-campaign contribution UI vs. Common Ground Budget setup.

Primary metric: total counted dollars from users who mark at least one project as “weak common-ground support.”

Hypothesis: Common Ground Budget materially increases funding from otherwise-lukewarm supporters.

Test two

No failure bonus vs. small failure bonus for early authorized contributors.

Primary metric: near-threshold conversion rate and early authorization rate.

Guardrail: sponsor-pool wastage from failed campaigns.

Test three

Pure capped QF bonus vs. diversity-adjusted capped QF bonus.

Primary metric: active cluster count per cleared project.

Guardrail: donor comprehension and complaint rate.

Test four




Manual user routing vs. optimizer-assisted default routing with user override.

Primary metric: payout throughput, threshold clears, and donor regret / override rate.

The point of the experiments is not to optimize engagement; it is to learn whether the mechanism actually turns cross-view weak support into more funded public goods without opening major gaming channels.

That is consistent with the site’s current measurement philosophy. 42

Rollout plan



Risks and limitations

The biggest strategic limitation is conceptual, not engineering. Forethought may simply be right that **voluntary causal mechanisms for non-excludable moral public goods are inherently hard**. If so, no web product will fully solve the free-rider problem. The proposed mechanism should therefore be judged against the right counterfactual: not “perfect public-goods provision,” but “more dollars moved from weak overlap into real cleared funding, with acceptable safety and governance costs.” ⁴³

The biggest mechanism risk is **gaming**. QF-style layers are vulnerable to Sybil attacks, collusion, and donor splitting, which is why real deployments treat identity and anti-fraud systems as part of the mechanism, not optional implementation detail. Moral Trade is already on the right track by tying allocation eligibility to identity confidence, using donor caps, and separating moral reputation from allocation power; the proposed design should tighten those practices, not relax them. ⁴⁴

The biggest moral risk is **mislabeling a contested good as “common ground.”** The current site already surfaces dissent signals and externality review, and that needs to become even more important once weak-support budgets are routable. The fix is procedural, not cosmetic: every project should have minimum-cluster breadth, visible dissent counts, and a hard rule that unresolved externality or anti-threat blockers prevent threshold clearance and capture. ⁴⁵

The biggest UX risk is **cognitive overload**. A mechanism that is theoretically elegant but hard to understand will fail in practice. Moral Trade’s current “choose your path,” worked examples, and plain-language primer are good assets here, and the new mechanism should preserve that style. The user-facing abstraction should be “set one common-ground budget,” not “configure a quadratic-assurance-allocation function.” ⁴⁶

The main operational limitation in this report is that much of Moral Trade is still public-pilot or demo-only, so some conclusions necessarily rest on theory plus product inspection rather than on a long empirical history of live rounds. I also did not identify every private or internal API behind the site, so some implementation details above are architecture proposals rather than confirmed internals. That uncertainty is one reason my confidence is moderate rather than extreme on exact parameter choices, even though the directional recommendation is clear. ⁴⁷

¹ ⁵ ⁸ ¹⁵ ²¹ <https://www.amirrorclear.net/files/moral-trade.pdf>
<https://www.amirrorclear.net/files/moral-trade.pdf>

² ⁹ <https://www.moraltrade.org/>
<https://www.moraltrade.org/>

³ ⁴ ⁶ ¹⁶ ¹⁷ ¹⁹ ²⁰ ²² ²³ ⁴³ <https://www.forethought.org/research/moral-public-goods-are-a-big-deal-for-whether-we-get-a-good-future>
<https://www.forethought.org/research/moral-public-goods-are-a-big-deal-for-whether-we-get-a-good-future>

⁷ <https://www.forethought.org/research/convergence-and-compromise>
<https://www.forethought.org/research/convergence-and-compromise>

¹⁰ <https://www.moraltrade.org/mpgf>
<https://www.moraltrade.org/mpgf>

¹¹ ²⁵ ⁴⁴ <https://www.moraltrade.org/mpgf/rounds/mpgf-assurance-round-demo-2026-05>
<https://www.moraltrade.org/mpgf/rounds/mpgf-assurance-round-demo-2026-05>

¹² ⁴⁵ <https://www.moraltrade.org/mpgf/pools>
<https://www.moraltrade.org/mpgf/pools>

¹³ ²⁸ <https://www.moraltrade.org/validation>
<https://www.moraltrade.org/validation>

14 24 <https://www.moraltrade.org/mpgf/metrics>

<https://www.moraltrade.org/mpgf/metrics>

18 29 <https://arxiv.org/abs/1809.06421>

<https://arxiv.org/abs/1809.06421>

26 27 <https://www.moraltrade.org/moral-trade/technical-spec>

<https://www.moraltrade.org/moral-trade/technical-spec>

30 <https://link.springer.com/content/pdf/10.1023/A%3A1004957109535.pdf>

<https://link.springer.com/content/pdf/10.1023/A%3A1004957109535.pdf>

31 <https://www.moraltrade.org/safety>

<https://www.moraltrade.org/safety>

32 <https://docs.stripe.com/identity>

<https://docs.stripe.com/identity>

33 <https://docs.stripe.com/payments/place-a-hold-on-a-payment-method>

<https://docs.stripe.com/payments/place-a-hold-on-a-payment-method>

34 <https://docs.stripe.com/payments/payment-intents>

<https://docs.stripe.com/payments/payment-intents>

35 https://docs.stripe.com/api/payment_intents/capture

https://docs.stripe.com/api/payment_intents/capture

36 37 38 39 <https://www.moraltrade.org/mpgf/governance>

<https://www.moraltrade.org/mpgf/governance>

40 41 42 <https://www.moraltrade.org/measurement>

<https://www.moraltrade.org/measurement>

46 <https://www.moraltrade.org/start>

<https://www.moraltrade.org/start>

47 <https://www.moraltrade.org/status>

<https://www.moraltrade.org/status>