

Improving MoralTrade.org for Real Moral Trade

Executive summary

MoralTrade.org is already unusually well aligned with the **spirit** of Toby Ord's proposal. It explicitly frames itself as a marketplace for "voluntary trade under disagreement," distinguishes formats such as pledge swaps and donation offsets, requires explicit terms and evidence rules, blocks threats and coercion, and keeps pilot statistics honest by separating worked examples from live activity. Those are real strengths, and they matter because Ord's core idea is not "charity matchmaking" in general; it is **trade made possible by differences in the parties' moral views**, constrained by voluntary agreement and judged relative to the no-trade default. ¹

The central weakness is not conceptual misalignment. It is that the site is still much closer to a **careful prototype** than to a functioning market. Publicly, the site reports **0 live offers, 8 worked examples, 2 public profiles, and \$0 reviewed offsets**. Matching is currently rule-based, evidence is largely manual, the platform disclaims escrow/custody/legal/tax services, and real-money public-goods checkout is still planned rather than active. In other words, the biggest problem is not that the idea is wrong; it is that the market is still too thin, too manually mediated, and too weakly institutionalized for users to rely on it at scale. ²

My highest-confidence recommendation is to **narrow the product** before widening it. The best near-term version of MoralTrade is not "an eBay for all moral services." It is a **verified coordination platform for bounded moral trades**, centered on two use cases:

- (1) **donation offsets into pre-approved compromise destinations / moral public goods**, and
- (2) **bounded pledge swaps with clear evidence rules and short time horizons**.

That focus best fits Ord's examples, Forethought's argument for moral public goods, and the current trust/compliance constraints of the product. ³

The next-most-important change is to build a **serious validator and evidence layer**. Ord identifies both ordinary compliance trust and "counterfactual trust" as major obstacles to realizing the gains from moral trade. The current site already leans in this direction with evidence-gated review, pending verification, staged privacy, and manual review queues; it should now formalize that into reviewer roles, evidence schemas, challenge windows, one-proof-one-claim controls, and visible completion states. Without that layer, the site risks looking thoughtful but not dependable. ⁴

The right governance model for the next phase is **centralized operator control with externally legible rules and third-party auditability**, not a DAO and not a fully decentralized registry. The site itself says it is "centralized for simplicity, portable later"; Forethought's recent writing likewise notes that centralized and decentralized implementations are both possible, and that decentralization mainly improves portability rather than obviously solving the core coordination problem. Early on, the binding constraints are safety review, legal compliance, dispute handling, and trust-building — all of which are easier under a centralized operating model with exportability than under tokenized or DAO governance. ⁵

Technically, MoralTrade should also resist the temptation to over-engineer. Given current public scale, the best fit is a **modular monolith** with a clean event log, not microservices and not an on-chain product. That is consistent with the “monolith first” / “microservice premium” software literature, with the site’s pilot stage, and with the fact that payment, identity, and fraud tooling are already much more mature in off-chain vendors such as Stripe and philanthropy infrastructure such as Every.org. ⁶

A concise prioritization is below.

Horizon	Priority move	Why it matters most now	Rough resource range
Short term	Narrow to verified donation offsets, public-goods contributions, and bounded pledge swaps	Maximizes Ord-alignment while minimizing legal, tax, escrow, and verification complexity	2–4 months; 3–5 core people; roughly \$120k–\$300k
Short term	Build validator workflows, evidence schemas, challenge windows, and completion states	Trust and counterfactuality are the core bottlenecks	Included above or + \$50k–\$150k
Medium term	Add tiered identity, reputation, reviewer marketplace, and API/webhook integrations	Converts prototype records into a reliable repeated marketplace	6–9 months; 5–8 people; roughly \$400k–\$900k
Long term	Add portable registry, third-party auditors, and threshold public-goods mechanisms such as dominant assurance contracts	Improves scale, portability, and public-goods liquidity once the core market is working	9–18 months; 7–12 people; roughly \$1.2M–\$3M

These estimates are my own rough implementation judgments, not sourced market prices.

My overall confidence is **high** that the right immediate move is to narrow scope and strengthen verification/governance, **moderate** that public-goods pooling should eventually use stronger mechanism design, and **moderate-to-low** on any recommendation that depends on unpublished facts such as the actual codebase, legal entity structure, target jurisdictions, or hidden user demand.

What the theory requires

Ord’s basic definition is crisp: **moral trade is trade made possible by differences in the parties’ moral views**. He distinguishes **pure moral trade**, where both parties see the result as morally better, from **mixed moral trade**, where at least one side also treats part of the gain as prudential. He also emphasizes that the relevant benchmark is the **default or status quo** — the no-trade outcome — and that the interesting region is the subset of options that is better than that default for both parties. He describes the gain above default as **moral surplus**. ⁷

That theoretical framing has direct product implications. A moral-trade platform should not merely collect preferences or list causes. It should force users to specify: the proposed action, the requested reciprocal action, the relevant baseline or default, the duration, the exit condition, and the evidence standard.

Otherwise the platform cannot tell whether a proposal is a real Pareto improvement over non-trade, or merely vague moral language. MoralTrade already asks for cause area, action, requested counterpart, expected impact, verification method, duration, payment cadence where relevant, and exit conditions; that is exactly the right design instinct. ⁸

Ord also argues that today's "market for moral trade is very incomplete," and explicitly imagines more sophisticated mechanisms — including market structures, bargaining, professionalization, and even a moral-trade version of eBay or Craigslist. That matters because MoralTrade should not measure success merely by elegance of examples or philosophical clarity. It should measure success by whether it is actually reducing market incompleteness: helping compatible but disconnected people discover each other, structure terms, verify performance, and repeat. ⁹

A second major theoretical implication is that **trust** is the hard problem. Ord distinguishes ordinary compliance/factual trust from **counterfactual trust**: whether the other party would have done the action anyway absent the trade. He notes that counterfactual trust is especially difficult to police, that longer horizons usually make it worse, and that shorter horizons and parties of known character help. He also notes that escrow-like arrangements can help with factual compliance even when specific performance cannot be forced. ¹⁰

That means a serious moral-trade product should prefer:

- **short, bounded commitments** over open-ended ones,
- **pre-registered baselines** over after-the-fact narratives,
- **reviewable evidence** over self-attestation,
- **repeat interaction and reputation** over one-shot anonymous deals, and
- **low-counterfactuality cases** first, because the more a trade depends on "what I would otherwise have done," the less objectively auditable it is. ¹¹

Forethought's work adds an important extension: not all gains from moral trade are bilateral. Some of the most important opportunities may be **moral public goods** — goods that many parties value for moral reasons, but that no one is individually motivated enough to fund on their own. Forethought explicitly links moral public goods to moral trade and argues that coordination mechanisms for such goods could be extremely important, especially when power is broadly distributed and when many agents share some consensus preferences. It also argues that voluntary mechanisms to fund moral public goods are promising but genuinely hard because of freeriding. ¹²

Forethought also introduces the main strategic danger: **threats and blockers**. In trade, both sides can agree the result is better than default. In threats, at least one side can view the world as worse simply because the threat exists. Forethought's convergence-and-compromise work therefore matters for site design: a moral-trade platform has to make anti-threat and anti-extortion checks a first-class product feature, not a policy footnote. ¹³

A compact synthesis is below.

Requirement implied by Ord and Forethought	Why it matters	What MoralTrade already has	What is still missing
Explicit no-trade baseline / default	Moral trade is evaluated against the status quo, not against abstract ideals	Some structured fields, duration, exit conditions, evidence rules	Stronger baseline declarations and counterfactual fields
Voluntary, non-threatening exchange	Trade is distinct from extortion or coercion	Strong safety language and blocked proposal classes	Better operational anti-threat review and appeals
Verifiable completion	Trust is a core bottleneck	Manual review, evidence rules, external payment evidence	Better validators, challenge windows, proof uniqueness, audit logs
Repeated-market trust and reputation	Thin one-shot markets are fragile	Profiles, ratings/ comments language exist in privacy docs, but little live usage	A real reputation model based on verified completion
Multi-sided liquidity	Moral trade markets are currently incomplete	Search, wish registry, rule-based suggestions	Market-making, better matching, incentives, narrow wedge focus
Support for moral public goods	Some gains are pooled, not bilateral	Public Goods Fund / MPGF conceptually present	Better mechanism design and stronger contribution accounting

This table is a synthesis of Ord's paper, Forethought's public-goods and threats arguments, and the current public product language. ¹⁴

Where MoralTrade.org stands now

Publicly, MoralTrade presents itself as a **prototype marketplace** for voluntary moral trade, donation offsets, and public-good coordination. It offers four main trade formats — **pledge swaps, donation offsets, paid action offers, and public-goods contributions** — plus a wish registry, public profiles, worked examples, and explanation pages on methodology, safety, privacy, and evidence standards. It also repeatedly explains that it **does not provide legal, tax, escrow, or custody services**. ¹⁵

The site is also admirably candid about pilot status. On the public homepage, it reports **0 live offers, 8 worked examples, 2 public profiles**, and **\$0 reviewed offsets**, and it explicitly says that worked examples are not live liquidity and that trust signals should not be invented before records exist. That is good product honesty, but it also means the site has not yet crossed the threshold from “coherent prototype” to “working exchange.” ¹⁶

The current design philosophy is unusually deliberate. Matching is described as **rule-based**, using cause areas, compatibility with payment or pledges, shared terms, and consent-gated previews “rather than AI inference.” The wish registry exposes only broad previews; exact wishes, constraints, identity details, and contact data remain behind privacy grants and consent stages. The platform is “centralized first, portable later,” with export/import language that gestures toward future interoperability without prematurely decentralizing the product. ¹⁷

That design has three real strategic virtues.

First, it is well aligned with the **anti-threat / anti-coercion** requirement. The site blocks proposals involving violence, illegal acts, fraud, extortion, doxxing, harassment, exploitation, or pressure on vulnerable people, and it avoids autonomous AI outreach, mass profile ingestion, and private-feed scraping. This is unusually important in a domain where “finding counterparties” can easily shade into surveillance or manipulation. ¹⁸

Second, it is well aligned with **privacy-by-design**. The privacy page clearly separates public profile data from private wish-profile data, limits source connections to consent scope and manual summaries, and describes field-level disclosure grants and portability. That is directionally consistent with ICO guidance that data protection should be built in “at the design phase and throughout its lifecycle,” and that data minimization means collecting only the minimum data necessary for a purpose. ¹⁹

Third, it makes serious progress on **legibility**. The reasoning/evidence standards page says a useful trade record must make the action, reciprocal request, verification rule, and safety boundary visible before anyone relies on the exchange. It requires public offers to show the action, reciprocal request, minimum reciprocal impact, duration, exit rules, evidence rules, and — for offsets — named compromise destination, unmatched-surplus rule, threshold, expiry, and anti-threat certification. That is an excellent foundation. ²⁰

The main product gaps are more operational than conceptual.

The biggest gap is **validator capacity**. Right now, “manual review required” appears everywhere, but the public site does not yet show a strong validator institution: no reviewer tiers, no SLAs, no conflict rules, no challenge window, no reliability scoring, and no public evidence-state taxonomy beyond pending/reviewed language. Because Ord’s trust problem is the bottleneck, this is the most important missing institution. ²¹

The second gap is **liquidity formation**. Search, rule-based matching, and worked examples are present, but there is not yet evidence of market-making tactics for overcoming the cold-start problem of a two-sided platform. Rochet and Tirole’s two-sided-market overview stresses that platforms have to get both sides “on board,” that pricing structure matters, and that residual cross-side externalities shape adoption. A moral-trade marketplace cannot assume that philosophical interest alone will create enough density to make matching easy. ²²

The third gap is **transactional infrastructure**. Publicly, the site currently uses three different patterns: a simple donation page that sends users to **Every.org** for a small set of verified funds; an MPGF contribution flow where users can submit **manual evidence of external payments** to approved destinations such as Open Collective or a fiscal host; and a “real-money MPGF contributions use Stripe Checkout after production readiness” plan that is explicitly not yet active. Those choices are prudent, but they also show that the money layer is still fragmented and provisional. ²³

A role-based view makes the current state clearer.



This flow is the right general shape for moral trade given Ord's emphasis on explicit default-improving exchange and trust, and it broadly matches the site's public language on drafting, matching, evidence review, and manual/public-goods workflows. ²⁴

Recommended product and governance model

The highest-leverage product decision is to **choose a narrow wedge**. My recommendation is:

Primary wedge:

verified donation offsets and public-goods contributions into pre-approved compromise destinations

Secondary wedge:

bounded pledge swaps with short horizons and strong evidence rules

Delay:

general paid-action offers as a mainstream feature

The reason is straightforward. Ord's own practical examples include offsetting opposed donations and short-horizon reciprocal commitments. Forethought's public-goods work suggests that shared compromise destinations can create gains that are easier to explain and sometimes easier to scale than bespoke bilateral moral barter. By contrast, paid action offers quickly move into employment/labor, regulated compensation, exploitation, and stronger AML/KYC burdens. ²⁵

The best near-term product is therefore not "everything for everyone," but a carefully staged ladder:

1. **Public-goods and donation-offset coordination** using verified external rails.
2. **Short, bounded pledge swaps** where evidence is legible.
3. **Only later**, category-constrained paid actions, once identity, dispute, and compliance workflows are mature. ²⁶

A comparison of core product modes is below.

Product mode	Ord alignment	Liquidity prospects	Trust burden	Legal / tax burden	Recommendation
Bilateral pledge swap	Very high	Moderate in niche communities	High counterfactual trust burden	Moderate	Keep, but constrain to short, bounded actions

Product mode	Ord alignment	Liquidity prospects	Trust burden	Legal / tax burden	Recommendation
Donation offset to compromise destination	Very high	High if destinations are pre-approved and legible	Lower than behavior-change swaps because receipts are easier to verify	Moderate-to-high	Make this a core wedge
Public Goods Fund / pooled compromise	High and strengthened by Forethought	Potentially high if seeded well	Medium	Moderate	Make this the central scaling path
Paid action offers	Mixed: compatible in principle, but drifts fastest toward ordinary labor-market and enforcement issues	Unclear	High	High	De-emphasize until later

This table is a design synthesis grounded in Ord's examples, Forethought's public-goods analysis, and the site's current product boundaries. ²⁷

What should change in the user journeys

For **donors**, MoralTrade should become more like Giving What We Can and Every.org in three respects: a strong history/dashboard layer, clearer use of third-party evaluative infrastructure, and explicit explanation of donation routing and discretion. GWWC shows donors their giving history and uses third-party evaluators and funding partners; Every.org routes the donation through Every.org, immediately issues a charitable receipt, and then disburses to the recommended nonprofit. MoralTrade does not need to copy those models exactly, but it should make the post-payment path much clearer than it currently is. ²⁸

For **sellers / offer creators**, the site should add stronger pre-publication onboarding. Its own reasoning standards say the next step should be "deeper onboarding and richer examples, not loosening these standards." I agree. High-integrity marketplaces do not become liquid by lowering the bar; they become liquid by making it easier to clear the bar. ²⁹

For **buyers / counterparties**, the site needs much stronger trust signals. eBay reduces trust barriers through buyer protection, feedback, rules against feedback manipulation, seller-protection policies, and category-specific authenticity services. MoralTrade cannot simply clone those institutions, but it needs equivalents: verified completion counts, dispute histories, reviewer confidence, evidence-grade labels, and standardized outcome states visible before a user commits. ³⁰

For **validators**, the current “manual review” concept should become an explicit product surface. Validators should have:

- a defined scope of review,
- evidence schemas by format,
- conflict-of-interest rules,
- a challenger / appeal lane,
- reviewer quality metrics, and
- a publishable reasoning summary for nontrivial disputes.

This is especially important because the current site already contemplates reviewer verification for the MPGF and admin review queues for reports, blocked wish profiles, payment-review requests, and disputes.

31

Governance choices

The right governance choice is not “trustless first.” It is **operator-legible first**.

Governance model	Advantages	Main liabilities	Best use
Centralized operator	Fast decisions, coherent policy, clear accountability, easier compliance	Single point of failure, weaker neutrality perceptions	Launch and early scale
Centralized operator plus third-party auditors / reviewers	Preserves speed while adding legitimacy and challengeability	More process overhead	Best next-step model
Federation / portable registry	Better portability and resilience, less lock-in	Harder policy consistency, harder abuse handling across nodes	Later, once standards stabilize
DAO / tokenized governance	Stronger ideological appeal to some users	High legal/governance complexity, Sybil and capture risks, weak fit for safety adjudication	Not recommended near term

This recommendation follows from MoralTrade’s own “centralized first, portable later” posture, Forethought’s observation that decentralization mainly improves portability rather than obviously solving the hard adoption problem, and the fact that launch-stage trust/safety/compliance problems are operational rather than constitutional. 5

The governance move I would make next is:

Central operator + published rulebook + independent reviewer panel + public quarterly transparency report.

That would borrow the best features of:

- GWWC’s transparency and explicit acknowledgement of conflicts, evaluators, and funding partners, ³²
- eBay’s explicit protection and policy structure, ³³
- and MoralTrade’s own existing emphasis on explicit boundaries and review before reliance. ³⁴

It would also make the platform more robust against Forethought’s threats concern, because “anti-threat certification” and coercion judgments become institutional processes rather than ad hoc moderator calls.

³⁵

Recommended technical architecture and operating metrics

Technical architecture

Publicly, MoralTrade does **not** specify its framework, database, hosting setup, or service decomposition. What it does specify is: rule-based matching, deterministic wish synthesis, privacy-gated data disclosure, centralized implementation, potential Stripe routing in some flows, and future export/import/schema portability. That is enough to make a high-confidence recommendation: **use a modular monolith with strong event logging** until real scale demands otherwise. ³⁶

The software-architecture literature is supportive here. Martin Fowler’s “Microservice Premium” argues that most systems should begin as a single monolithic application with good modularity, and his “Monolith First” discussion treats that as the default position for many systems, even while acknowledging situations where early decomposition helps. His microservices guide also emphasizes that microservices are independently deployable small services with nontrivial operational consequences. At MoralTrade’s present public scale, paying the microservice premium would be premature. ³⁷

A comparison follows.

Stack option	Fit for current stage	Strengths	Weaknesses	Recommendation
Modular monolith + Postgres + background jobs + event log	Excellent	Simpler ops, faster iteration, easier integrity constraints, easier reviewer tooling	Less independent scaling	Recommended now
Event-driven modular monolith with extractable domains	Very good	Preserves future portability, supports webhooks and audit trails	Slightly more engineering discipline required	Recommended target shape

Stack option	Fit for current stage	Strengths	Weaknesses	Recommendation
Full microservices	Poor-to-moderate	Independent scaling, team autonomy	High complexity, observability overhead, distributed-failure modes	Delay
On-chain / blockchain-first	Poor	Portability, public verifiability in narrow cases	User friction, privacy conflicts, custody/regulatory burden, weak fit for nuanced disputes	Avoid for launch

This is a stage-of-company judgment based on present public product maturity and the cited architecture references. ³⁸

A sensible near-term architecture would include these modules inside one deployable system:

- identity and role management,
- offer drafting and taxonomy,
- matching and notification,
- evidence storage and review workflow,
- dispute handling,
- public-goods pool accounting,
- and analytics / fraud / audit logging.

That should be paired with **external specialists** for money and identity rather than rebuilding them in-house. Stripe Connect is explicitly built for platforms and marketplaces that move money among multiple parties; its Identity product supports document, selfie, ID number, address, and phone checks; and its Radar products provide rules, risk scoring, analytics, and merchant-risk signals. Those are much stronger primitives than most early-stage teams can replicate. ³⁹

For public-goods donation infrastructure, Every.org is an attractive complement because it already exposes a Charity API, supports a wide range of donation methods, and issues receipts as a 501(c)(3) intermediary before disbursing to the recommended nonprofit. MoralTrade's current Every.org donation usage is therefore pointing in the right direction: use specialist rails wherever possible rather than trying to become a fully fledged donations processor immediately. ⁴⁰

UX, accessibility, privacy, and internationalization

The most important UX change is **better onboarding, not more browsing chrome**. The site already says the next step is deeper onboarding and richer examples. The onboarding should become role-specific:

- "I want to donate and find a credible compromise destination."
- "I want to offer a bounded action for a reciprocal commitment."

- “I want to validate evidence.”
- “I want to browse only high-confidence, reviewable offers.” 29

Trust signals should become more explicit and more structured. eBay teaches three especially useful lessons here:

- protection and dispute rules should be obvious before transacting, 41
- reputation has to be protected against manipulation, 42
- and category-specific verification can matter more than generic ratings in high-trust categories. 43

For MoralTrade, the analogue is not “star ratings everywhere.” It is:

- verified completion count,
- evidence-grade badge,
- current review state,
- reviewer-confidence band,
- dispute history,
- and “same proof already used elsewhere?” flags.

Accessibility should be treated as a launch-quality requirement, not a cleanup task. W3C’s WCAG 2.2 is the current international standard family for making web content accessible, and W3C explicitly encourages using the latest version. For a site centered on complex forms, review states, and consent gates, this has concrete implications: keyboard navigation, visible focus management, clear step states, color-independent status cues, readable evidence instructions, and error messages that identify missing fields precisely. 44

Internationalization also matters more than it might seem. W3C states that web standards are optimized for interoperability, accessibility, privacy, security, and internationalization, and its internationalization guidance emphasizes language-aware content and broad deployment support. Because moral trade depends so heavily on semantic nuance, any future multilingual rollout needs careful language tagging, late localization, and a controlled vocabulary for cause areas, evidence types, and review states. 45

Privacy should remain a product differentiator. The current semi-private, consent-gated design is directionally strong, and it fits ICO guidance on privacy/data protection by design and data minimization. I would strengthen, not weaken, this posture:

- pseudonymous public handles by default,
- field-level grants,
- minimal data retention by feature,
- separate access for validator teams versus matchmaking teams,
- and differential disclosure for broad preview, specific terms, and contact exchange. 46

Incentives and economic design

MoralTrade is a **two-sided platform**. Rochet and Tirole’s framework matters here: platforms must get both sides on board, and price structure matters because there are residual externalities between sides. In practice, that means MoralTrade should not expect to fund itself early by extracting fees from every participant equally. One side usually has to be subsidized. 22

My recommendation is:

- **No listing fees** in pilot and early scale.
- **No success fee** on public-goods or donation-offset flows that use nonprofit rails.
- **Philanthropic subsidy** for validator operations and early review.
- If needed later, charge **small fees only on professional / institutional workflows**, not on ordinary individual trades.
- Consider **review bounties** rather than blanket moderator labor for complex cases.

That is closer to Every.org's and GWWC's ethos than to eBay's transaction-fee model, and I think that is the right fit here. Every.org's platform emphasizes zero platform fees and optional support donations; GWWC emphasizes impact infrastructure, recommendations based on third-party evaluators, and donor dashboards rather than extracting marketplace rents from every interaction. ⁴⁷

For the **Public Goods Fund**, the most promising later-stage mechanism is a **dominant assurance contract** rather than immediate quadratic funding. Tabarrok's dominant assurance-contract result is attractive here because it directly addresses first-mover reluctance in threshold public-goods situations by adding a refund bonus if the threshold is not met, making contribution a dominant strategy in the modeled setting. That is simpler, closer to the current MPGF "pledge rehearsal" concept, and easier to explain to users than more sophisticated matching mechanisms. ⁴⁸

Quadratic-funding-style mechanisms may eventually be useful for moral public goods, and the original papers argue they can improve public-goods provision by amplifying broad support. But I would treat them as a **long-run option only after** strong identity, anti-Sybil controls, and serious auditing are in place. ⁴⁹

Measurement and evaluation

The best KPI for a marketplace is not traffic. It is **credible liquidity**.

For MoralTrade, I would define five KPI families:

KPI family	What to measure	Why it matters
Market liquidity	active verified offers, median time to first credible match, match rate, repeat users	Tells you whether the market exists
Trust and verification	evidence pass rate, dispute rate, reviewer turnaround, appeal overturn rate	Tells you whether users can rely on outcomes
Impact / coordination	verified redirected dollars, verified reciprocal actions, public-goods dollars unlocked, matched vs unmatched surplus	Tells you whether the platform coordinates anything meaningful
Safety	threat/coercion blocks, abuse reports, false-positive review rate	Tells you whether the anti-threat posture is working

KPI family	What to measure	Why it matters
Fraud / integrity	duplicate-proof rate, suspicious cluster rate, chargebacks if payments go live, account-risk flags	Tells you whether the market can be gamed

These metrics are strongly supported by the reputation-systems literature, Stripe’s risk tooling, and the site’s own emphasis on review-before-reliance and explicit pilot statistics. Reputation systems are central trust-building mechanisms in online interactions, but they are vulnerable to attack and manipulation, which means a serious marketplace must log and analyze integrity failures, not just completions. ⁵⁰


The specific impact metric I would add — because this product is philosophically distinctive — is a **counterfactual-confidence score**. This should not pretend to solve Ord’s counterfactual-trust problem. It should merely make the uncertainty visible by combining:

- pre-trade baseline declaration,
- participant self-report,
- counterparty assessment,
- reviewer confidence,
- and whether independent timestamps/evidence support the claim.

Ord is explicit that counterfactual trust is hard and often not fully policeable. The right product response is not false precision; it is structured uncertainty. ⁵¹

Roadmap, risks, and limitations

Suggested roadmap



The dates are illustrative and reflect my own implementation judgment.

A practical milestone plan is below.

Phase	Main deliverables	Staffing	Rough cost
Short term	Reposition product around verified offsets / public goods / bounded pledge swaps; role-specific onboarding; evidence templates; reviewer console; proof uniqueness checks; public status states	1 product lead, 2 full-stack engineers, 1 design/product-ops hybrid, 0.25 legal/compliance, part-time reviewer lead	\$120k–\$300k
Medium term	Identity tiers, verification provider integration, reputation model based on verified completion, reviewer marketplace, analytics dashboards, partner destination registry	+1 backend/data engineer, +1 ops/reviewer manager, +0.5 trust & safety	\$400k–\$900k
Long term	Portable profiles / APIs, third-party audit process, public-goods threshold mechanisms, federation experiments, institutional accounts	7–12 total people depending on scope	\$1.2M–\$3M

Main risks and mitigations

The biggest **legal/tax** risk is pretending that all reciprocal moral-trade donation flows are ordinary charitable gifts. They are not automatically so. IRS guidance distinguishes **quid pro quo contributions**, where the donor receives goods or services in return, and requires disclosure in various cases; it also stresses that donor-advised-fund sponsors must retain legal control over contributed assets. UK Gift Aid rules similarly require proper charity status/conditions, and not all payments qualify. Because MoralTrade's central mechanic often links a donation to reciprocal consideration, tax treatment is plainly fact-specific. The platform is right to disclaim legal/tax advice; it should go further and make potentially non-deductible flows visually distinct unless structured through approved partners under explicit rules. ⁵²

The second legal boundary concerns **political contributions**. This is one place where the current site is especially wise: it blocks political campaign contribution offsets, and Ord himself notes that Repledge-style political donation markets ran into unresolved legal questions. That boundary should remain hard. ⁵³

The biggest **ethical** risk is that the site could unintentionally legitimize coercion disguised as trade. Forethought's threats argument and the site's own anti-threat rules point in the same direction: anti-threat certification, coercion review, and vulnerable-user protections are not optional. They are part of the product definition. ⁵⁴

The biggest **security/integrity** risks are Sybil behavior, duplicate or fake evidence, manipulated ratings, and marketplace gaming. The reputation-systems literature is clear that these systems are attack surfaces, not just UX layers; OWASP ASVS provides a framework for testing application security controls; and Stripe's risk stack shows why merchant- and payment-level screening should be externalized when possible. MoralTrade should treat security and marketplace integrity as one combined discipline. ⁵⁵

The biggest **scalability** risk is thin-market failure. Two-sided platforms live or die on cross-side liquidity, and Ord's "incomplete market" problem is not solved by publishing elegant examples. The antidote is focus: a narrow wedge, market-making on that wedge, subsidized early review, and strong repeat-use loops. 56

Open questions and limitations

This assessment is necessarily limited by what is public. Several important specifics are **unspecified** on public pages:

- the actual backend architecture, framework, and database,
- legal entity structure and target jurisdictions,
- current user research and internal analytics,
- reviewer staffing and moderation capacity,
- payments-provider account configuration,
- and whether there is meaningful off-platform demand already waiting to be activated.

Because of that, I have high confidence in the **direction** of the recommendations, but only moderate confidence in later-stage mechanism and governance choices that depend on those hidden facts.

The core conclusion, however, is robust:

MoralTrade.org should not try to become a broad moral-services bazaar yet. It should become the most trustworthy place on the internet to structure, verify, and record bounded moral trades — especially verified donation offsets and moral-public-goods coordination. That path is the strongest fit to Ord's theory, to Forethought's extensions, to the current site's own design language, and to what existing online marketplaces have learned about trust, liquidity, evidence, review, and policy clarity. 57

1 3 4 7 8 9 10 11 14 24 25 27 51 57 <https://www.amirrorclear.net/files/moral-trade.pdf>
<https://www.amirrorclear.net/files/moral-trade.pdf>

2 15 16 21 38 <https://www.moraltrade.org/>
<https://www.moraltrade.org/>

5 17 36 <https://www.moraltrade.org/methodology>
<https://www.moraltrade.org/methodology>

6 37 <https://martinfowler.com/bliki/MicroservicePremium.html>
<https://martinfowler.com/bliki/MicroservicePremium.html>

12 <https://www.forethought.org/research/moral-public-goods-are-a-big-deal-for-whether-we-get-a-good-future>
<https://www.forethought.org/research/moral-public-goods-are-a-big-deal-for-whether-we-get-a-good-future>

13 35 54 <https://www.forethought.org/research/convergence-and-compromise>
<https://www.forethought.org/research/convergence-and-compromise>

18 34 <https://www.moraltrade.org/safety>
<https://www.moraltrade.org/safety>

19 46 <https://www.moraltrade.org/privacy>

<https://www.moraltrade.org/privacy>

20 29 <https://www.moraltrade.org/reasoning-standards>

<https://www.moraltrade.org/reasoning-standards>

22 56 https://web.mit.edu/14.271/www/rochet_tirole.pdf

https://web.mit.edu/14.271/www/rochet_tirole.pdf

23 <https://www.moraltrade.org/donate>

<https://www.moraltrade.org/donate>

26 31 <https://www.moraltrade.org/mpgf/contribute>

<https://www.moraltrade.org/mpgf/contribute>

28 <https://www.givingwhatwecan.org/pledge>

<https://www.givingwhatwecan.org/pledge>

30 33 41 <https://www.ebay.com/help/policies/ebay-money-back-guarantee-policy/ebay-money-back-guarantee-policy?id=4210>

<https://www.ebay.com/help/policies/ebay-money-back-guarantee-policy/ebay-money-back-guarantee-policy?id=4210>

32 <https://www.givingwhatwecan.org/transparency>

<https://www.givingwhatwecan.org/transparency>

39 <https://docs.stripe.com/connect>

<https://docs.stripe.com/connect>

40 <https://www.every.org/charity-api>

<https://www.every.org/charity-api>

42 <https://www.ebay.com/help/buying/leaving-feedback-sellers/leaving-feedback-sellers?id=4007>

<https://www.ebay.com/help/buying/leaving-feedback-sellers/leaving-feedback-sellers?id=4007>

43 <https://www.ebay.com/help/selling/selling-tools/ebay-authenticity-guarantee?id=4644>

<https://www.ebay.com/help/selling/selling-tools/ebay-authenticity-guarantee?id=4644>

44 <https://www.w3.org/WAI/standards-guidelines/wcag/>

<https://www.w3.org/WAI/standards-guidelines/wcag/>

45 <https://www.w3.org/standards/>

<https://www.w3.org/standards/>

47 <https://www.every.org/pricing>

<https://www.every.org/pricing>

48 https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4682

https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4682

49 <https://www.microsoft.com/en-us/research/publication/liberal-radicalism-a-flexible-design-for-philanthropic-matching-funds/>

<https://www.microsoft.com/en-us/research/publication/liberal-radicalism-a-flexible-design-for-philanthropic-matching-funds/>

50 <https://presnick.people.si.umich.edu/papers/cacm00/reputations.pdf>

<https://presnick.people.si.umich.edu/papers/cacm00/reputations.pdf>

⁵² <https://www.irs.gov/charities-non-profits/substantiating-charitable-contributions>
<https://www.irs.gov/charities-non-profits/substantiating-charitable-contributions>

⁵³ <https://www.moraltrade.org/faq>
<https://www.moraltrade.org/faq>

⁵⁵ https://www.ime.usp.br/~cris/aulas/17_2_6906/projetos/Reputation-systems-ACMCS09.pdf
https://www.ime.usp.br/~cris/aulas/17_2_6906/projetos/Reputation-systems-ACMCS09.pdf